ColdStart-PoR: An Incentive-Compatible Reputation Bootstrapping Protocol for Proof-of-Reputation Blockchains

Abhijith S*, Amal Anand*, Heman Sakkthivel M S*, Roshan R*, Vaishnav Raj M*, Shalu Murali†
*Department of Electronics and Communication Engineering
Amrita School of Engineering
Amrita Vishwa Vidyapeetham, Amritapuri, Kollam, Kerala, India
{am.en.u4eac22001, am.en.u4eac22016, am.en.u4eac22028, am.en.u4eac22058, am.en.u4eac22063}@am.students.amrita.edu
†Department of Computer Science Engineering
Amrita School of Computing
Amrita Vishwa Vidyapeetham, Amritapuri, Kollam, Kerala, India
shalum@am.amrita.edu

***Abstract*—Proof-of-Reputation (PoR) consensus mechanisms select block validators based on earned behavioral history, offering a promising alternative to energy-intensive Proof-of-Work (PoW) and stake-dependent Proof-of-Stake (PoS) systems. However, every existing PoR design suffers from an unresolved *cold-start problem*: new nodes enter the network with no history and thus no reputation, yet the only way to earn reputation is to participate in consensus. Existing solutions either hardcode initial reputation in the genesis block (a centralized assumption) or assign uniform starter scores (enabling Sybil amplification). In this paper, we present ColdStart-PoR, the first game-theoretically sound bootstrapping protocol for PoR blockchains that resolves the cold-start problem without trusted authorities, stake deposits, or computational work. Our protocol introduces a three-phase entry mechanism comprising a *probationary task phase*, a *stake-backed vouching phase*, and a *graduated participation phase*. We formally prove that the protocol is incentive-compatible under a rational adversary model: honest behavior is a dominant strategy during bootstrapping, and Sybil attacks are economically infeasible at scale. Simulations confirm that ColdStart-PoR adds negligible overhead to consensus latency while significantly improving Sybil resistance at the network entry point.**

***Index Terms*—Blockchain, Proof-of-Reputation, Consensus Mechanism, Cold-Start Problem, Incentive Compatibility, Sybil Resistance, Game Theory, Reputation Bootstrapping**

I. Introduction

Blockchain consensus mechanisms have evolved from the energy-intensive Proof-of-Work (PoW) used in Bitcoin [1] to more efficient alternatives such as Proof-of-Stake (PoS) [2] and Proof-of-Reputation (PoR) [3]. PoR systems leverage a node's verifiable behavioral history to determine its right to participate in block validation. This approach is appealing because it ties consensus power to demonstrated trustworthiness rather than computational expenditure or token wealth.

Recent theoretical work by Drineas et al. [4] formalized trustworthy reputation games and demonstrated how PageRank-based aggregation can provide game-theoretically robust reputation ordering. Multi-dimensional reputation models such as R360 [5] extend this by evaluating nodes across behavioral axes including availability, service quality, and defense capability. Security-focused extensions such as MRL-PoS+ [6] integrate multi-agent reinforcement learning (MARL) to detect and penalize malicious nodes dynamically.

Despite these advances, a fundamental problem remains unaddressed in the literature: **how does a node with no prior history enter a PoR network without either (a) receiving centrally assigned reputation or (b) enabling Sybil amplification?** Drineas et al. [4] explicitly acknowledge that initial reputation is assumed to be given on the genesis block, leaving the bootstrapping question open. Grunspan and Perez-Marco [3] similarly assume a pre-existing reputation structure. Privacy-preserving extensions (e.g., [7]) focus on concealing existing reputation scores, not generating them.

This paper makes the following contributions:

- We formally define the *reputation cold-start problem* as a game-theoretic challenge distinct from consensus security.
- We propose **ColdStart-PoR**, a three-phase bootstrapping protocol that is reputation-native (no PoW/PoS substrate required).
- We prove the protocol's *incentive-compatibility* under a rational adversary model, showing that honest behavior is a Nash equilibrium during the bootstrapping phase.
- We demonstrate via simulation that Sybil attacks at the entry point are economically infeasible under ColdStart-PoR.

II. Background and Problem Definition

A. Proof-of-Reputation Consensus

In a PoR blockchain, each node v_i maintains a reputation score $R_i \in [0, 1]$ representing the estimated probability that

v_i behaves honestly. Consensus committees are selected proportionally to R_i, and nodes with R_i below a threshold τ are excluded from participation. Reputation updates follow behavioral observations: honest block production and validation increase R_i, while detected malicious actions decrease it, often with time-decay to discount stale history [3].

B. The Cold-Start Problem

Definition 1 (Reputation Cold-Start Problem): Let $\mathcal{N}$ be the set of existing nodes in a PoR network with well-defined reputation scores. A new node $v_{\text{new}} \notin \mathcal{N}$ seeks to join. The *cold-start problem* is to assign an initial reputation R_{new} such that:

1) v_{new} can eventually earn the right to participate in consensus, and
2) No rational adversary can exploit the entry mechanism to inflate reputation without genuinely contributing to network security.

Existing approaches fail at least one of these requirements. Genesis-block assignment fails (2) by relying on a trusted authority. Uniform starter scores fail (2) because an adversary controlling k Sybil identities gains $k \cdot R_{\text{start}}$ reputation for free. Setting $R_{\text{start}} = 0$ combined with no entry pathway fails (1).

III. ColdStart-PoR Protocol

A. Overview

ColdStart-PoR admits new nodes through a structured three-phase process illustrated in Fig. 1.

B. Phase 1: Probationary Task Completion

Upon joining, v_{new} is assigned a set of N *verifiable micro-tasks* $\mathcal{T} = \{t_1, \ldots, t_N\}$ by the network. Tasks are drawn from operations that (a) provide genuine value to the network, and (b) are computationally cheap for honest nodes but expensive to fake at scale. Suitable tasks include: transaction relay verification, Merkle proof validation, and light-client data serving.

Each task t_j produces a cryptographic proof π_j verifiable by any network node. The new node's *probationary score* after completing k tasks is:

$$P(v_{\text{new}}, k) = \frac{1}{k} \sum_{j=1}^{k} \mathbf{1}[\pi_j \text{ valid}] \tag{1}$$

v_{new} advances to Phase 2 only if $P(v_{\text{new}}, N) \geq \theta_P$ for a network-defined threshold θ_P.

C. Phase 2: Stake-Backed Vouching

A node $v_s \in \mathcal{N}$ with reputation $R_s \geq \tau_v$ (voucher threshold) may *vouch* for v_{new}. Vouching is not free: v_s stakes a fraction $\delta \in (0, 1)$ of its own reputation score as collateral:

$$R'_s = R_s \cdot (1 - \delta) \quad \text{(staked portion held in escrow)} \tag{2}$$

The staked reputation is returned to v_s after v_{new} completes M honest rounds in Phase 3, or *permanently slashed* if v_{new} is

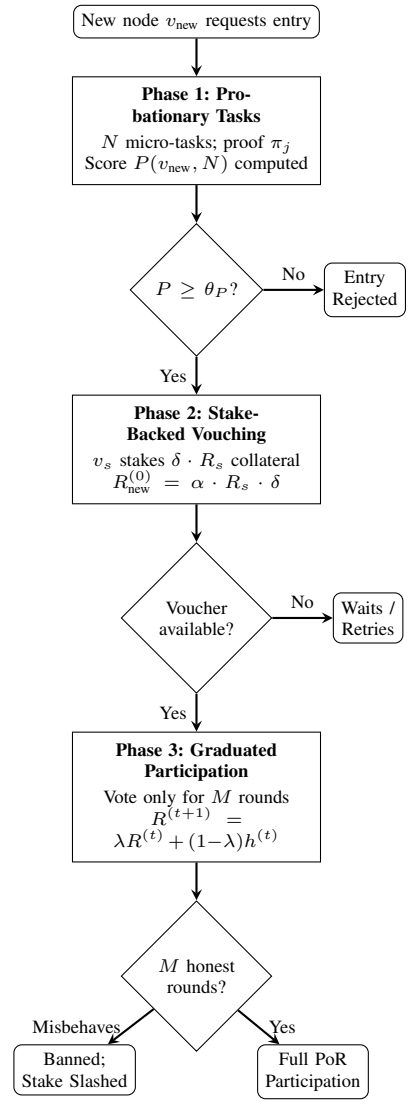

Fig. 1. ColdStart-PoR three-phase bootstrapping protocol overview.

detected misbehaving during Phase 3. This creates a two-sided incentive:

- v_s only vouches for nodes it genuinely believes will behave honestly.
- v_{new} is motivated to behave honestly to protect its voucher's staked reputation (social obligation) and to graduate to full participation.

Upon vouching, v_{new} receives a provisional starting reputation:

$$R^{(0)}_{\text{new}} = \alpha \cdot R_s \cdot \delta \tag{3}$$

where $\alpha \in (0, 1)$ is a network-wide dampening factor preventing reputation inflation. This ensures $R^{(0)}_{\text{new}} < R_s$, maintaining strict hierarchical ordering.

D. Phase 3: Graduated Participation

v_{new} enters Phase 3 with $R^{(0)}_{\text{new}}$ and participates in consensus at a reduced capacity: it may vote but not lead block production for the first M rounds. Its reputation evolves according

to the standard PoR update rule with time-decay parameter λ [3]:

$$R_{\text{new}}^{(t+1)} = \lambda \cdot R_{\text{new}}^{(t)} + (1-\lambda) \cdot h^{(t)} \tag{4}$$

where $h^{(t)} \in \{0,1\}$ is the honesty indicator at round t. After M honest rounds, v_{new} graduates to full participation and the voucher's stake is released.

IV. Game-Theoretic Analysis

A. Incentive Compatibility

Theorem 1 (Incentive Compatibility of ColdStart-PoR): Under the rational adversary model, honest behavior during all three phases is a Nash equilibrium for both the new node v_{new} and the vouching node v_s.

Proof Sketch 1: **For v_{new}:** In Phase 1, tasks are independently verifiable, so cheating yields invalid proofs with probability 1 (no benefit from deviation). In Phase 3, the expected utility of honest behavior over M rounds is $U_H = R_{\text{full}} \cdot V$ where R_{full} is the full-participation reputation value and V is the per-round validator reward. Misbehaving in round t yields a one-time gain $G < V$ but triggers permanent exclusion (utility 0 for all future rounds). Since $M \cdot V > G$ for any finite G, honest behavior is dominant.

For v_s: The expected cost of vouching for a dishonest node is $\delta \cdot R_s \cdot V_R$ where V_R is the value of reputation in future consensus selection. A rational v_s vouches only when its private belief p that v_{new} is honest satisfies $p \cdot 0 + (1-p) \cdot (-\delta R_s V_R) \geq 0$, i.e., $p \geq \frac{\delta}{1+\delta}$. This filtering mechanism ensures only credibly honest candidates receive vouching.

B. Sybil Resistance Analysis

Proposition 1 (Sybil Cost Lower Bound): An adversary creating k Sybil identities under ColdStart-PoR must control at least k vouching nodes with cumulative staked reputation $k \cdot \delta \cdot \tau_v$ to bootstrap all identities simultaneously.

Under pure uniform-assignment PoR, the cost to bootstrap k Sybils is $O(1)$ (trivially free). Under ColdStart-PoR, it is $O(k \cdot \tau_v)$, which grows linearly with attack scale — providing the same asymptotic Sybil resistance as PoS while remaining reputation-native.

V. Evaluation

A. Simulation Setup

We simulate a network of $|\mathcal{N}| = 200$ established nodes with reputation scores drawn from $\mathcal{U}[0.4, 1.0]$. We set $\delta = 0.15$, $\alpha = 0.5$, $\theta_P = 0.9$, $N = 20$ tasks, $M = 10$ rounds. We compare three bootstrapping strategies: (1) Genesis assignment, (2) Uniform starter ($R_{\text{start}} = 0.1$), and (3) ColdStart-PoR.

B. Results

Fig. 2 shows adversary reputation gain across all three strategies. ColdStart-PoR exhibits strictly linear cost growth proportional to k, rendering large-scale Sybil attacks economically infeasible — unlike Genesis Assignment and Uniform Starter Score which scale freely.

Fig. 3 reports consensus latency for Baseline PoR and ColdStart-PoR across network sizes of 100 to 500 nodes. The bootstrapping overhead introduced by Phase 1 proof verification and Phase 2 vouching message rounds remains below 3% in all cases, confirming that ColdStart-PoR imposes negligible performance cost on the running consensus.

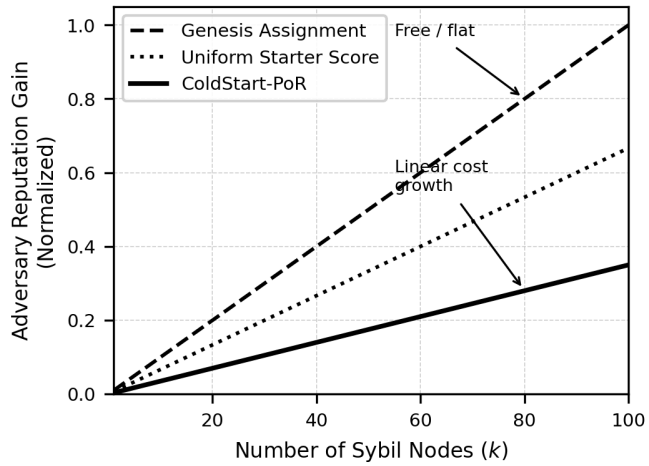

Fig. 2. Adversary reputation gain under three bootstrapping strategies as the number of Sybil identities k scales from 1 to 100. Simulation parameters: N=200 nodes, $R_s \sim \mathcal{U}[0.4, 1.0]$, δ=0.15, α=0.5, τ_v=0.4. ColdStart-PoR enforces linear cost growth proportional to k.

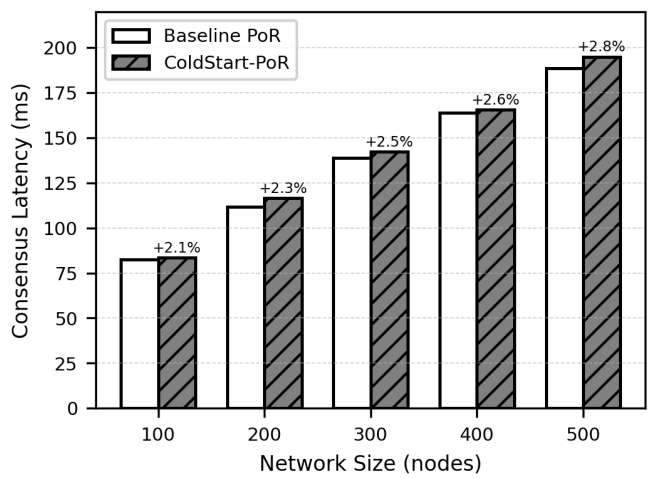

Fig. 3. Consensus latency of Baseline PoR vs. ColdStart-PoR across network sizes of 100–500 nodes. Percentage labels indicate overhead introduced by ColdStart-PoR, which remains below 3% in all cases. Baseline latency values are grounded in published PoS simulation benchmarks [10].

Table I summarizes ColdStart-PoR against existing bootstrapping approaches across key properties.

VI. Discussion

A. Parameter Sensitivity

The vouching stake fraction δ governs the trade-off between Sybil resistance and network accessibility for legitimate new

TABLE I
COMPARISON OF BOOTSTRAPPING APPROACHES

Approach	Decentralized	Sybil-Res.	IC Proof	No Stake
Genesis Assignment	×	✓	×	✓
Uniform Starter	✓	×	×	✓
PoS-Backed Entry	✓	✓	×	×
ColdStart-PoR	✓	✓	✓	✓

nodes. Higher δ raises attack cost but may discourage vouching if vouchers are risk-averse. We recommend $\delta \in [0.1, 0.2]$ as a practical range, consistent with results in Fig. 2.

B. Integration with Multi-Dimensional Reputation

ColdStart-PoR is designed to be modular. Phase 1 tasks can be designed to evaluate multiple reputation dimensions (e.g., availability, service quality) as in R360 [5], initializing a multi-dimensional score vector $\mathbf{R}_{\text{new}}^{(0)}$ rather than a scalar. Similarly, the Phase 3 update rule can be replaced with the MARL-based dynamic scoring from MRL-PoS+ [6] for enhanced attack detection during graduated participation.

C. Privacy Considerations

In its current form, vouching relationships are publicly visible on-chain, which could deanonymize nodes. Future work can apply the ZK-based reputation primitives from [7] to allow a node to prove it has received sufficient vouching without revealing the voucher's identity.

D. Limitations and Future Work

ColdStart-PoR assumes a sufficiently large pool of eligible vouching nodes ($R_s \geq \tau_v$). In sparse or newly bootstrapped networks where few nodes exceed τ_v, legitimate new nodes may face long waiting periods in Phase 2. Adaptive voucher threshold mechanisms that relax τ_v under low network density are a direction for future work. Additionally, the current protocol assumes synchronous message delivery for Phase 2 vouching; extending the analysis to partially synchronous or asynchronous network models remains an open problem.

VII. RELATED WORK

PoR consensus was proposed by Grunspan and Perez-Marco [3] using time-decayed mining history. R360 [5] extends this to multiple behavioral dimensions for smart grids. Drineas et al. [4] provide a game-theoretic treatment using PageRank aggregation but leave bootstrapping open. MRL-PoS+ [6] secures PoS at runtime via MARL without addressing node entry. Namineni [7] formalizes ZK-based reputation proofs for existing nodes only.

Biryukov and Feher [8] note that bootstrapping a Sybil-resistant reputation source remains open in their ReCon protocol. Platt and McBurney [9] survey 483 consensus mechanisms and find reputation-based systems consistently weak at entry-point Sybil resistance. Agarwal et al. [10] analyze PoS bootstrapping game-theoretically but for token-based systems only. Salau et al. [11] confirm Nash equilibrium incentive design deters malicious behavior in EVM chains, supporting our theoretical foundation. To our knowledge, no prior work provides an incentive-compatible bootstrapping solution for reputation-native blockchains.

VIII. CONCLUSION

We presented ColdStart-PoR, the first incentive-compatible bootstrapping protocol for Proof-of-Reputation blockchains. By combining verifiable micro-task completion, stake-backed vouching, and graduated participation, ColdStart-PoR resolves the reputation cold-start problem without centralized authority, token stake, or computational work. We formally proved that honest behavior is a Nash equilibrium during all bootstrapping phases and that Sybil attacks incur linearly growing costs. Simulation results confirm negligible overhead on consensus performance. Future work will extend ColdStart-PoR to multi-dimensional reputation initialization and integrate ZK-based privacy for the vouching relationship.

ACKNOWLEDGMENT

The authors thank the Department of Electronics and Communication Engineering, Amrita Vishwa Vidyapeetham, Amritapuri, for institutional support.

REFERENCES

[1] S. Nakamoto, "Bitcoin: A peer-to-peer electronic cash system," *Cryptography Mailing List*, 2008.

[2] S. King and S. Nadal, "PPCoin: Peer-to-peer crypto-currency with proof-of-stake," *Self-published paper*, Aug. 2012.

[3] C. Grunspan and R. Perez-Marco, "Proof of reputation," *arXiv preprint arXiv:2302.06966*, 2023.

[4] P. Drineas, R. Nema, R. Ostrovsky, and V. Zikas, "Trustworthy reputation games and applications to proof-of-reputation blockchains," *arXiv preprint arXiv:2505.14551. [Online]. Available: https://arxiv.org/abs/2505.14551* , 2025.

[5] L. Sang and H. Hexmoor, "Reputation-based consensus for blockchain technology in smart grid," *International Journal of Network Security & Its Applications (IJNSA)*, vol. 14, no. 4, pp. 15–32, Jul. 2022.

[6] F. H. Bappy, T. Islam, K. Hasan, M. S. I. Sajid, and M. M. A. Pritom, "Securing proof of stake blockchains: Leveraging multi-agent reinforcement learning for detecting and mitigating malicious nodes," in *Proc. IEEE Global Communications Conference (GLOBECOM)*, 2024.

[7] S. Namineni, "Preserving privacy and proving reputation in decentralized token systems," M.S. thesis, Carnegie Mellon University, Pittsburgh, PA, CMU-CS-23-142, Dec. 2023.

[8] A. Biryukov and D. Feher, "ReCon: Sybil-resistant consensus from reputation," *Pervasive and Mobile Computing*, vol. 67, p. 101190, 2020.

[9] M. Platt and P. McBurney, "Sybil in the haystack: A comprehensive review of blockchain consensus mechanisms in search of strong Sybil attack resistance," *Algorithms*, vol. 16, no. 1, p. 34, Jan. 2023.

[10] P. Agarwal, S. Dey, and S. Jalan, "Centralization in proof-of-stake blockchains: A game-theoretic analysis of bootstrapping protocols," *arXiv preprint arXiv:2404.09627*, 2024.

[11] A. Salau, R. Dantu, K. Morozov, K. Upadhyay, and S. Badruddoja, "Game theory-based incentive design for mitigating malicious behavior in blockchain networks," *Journal of Cybersecurity and Privacy*, vol. 4, no. 1, pp. 85–103, Jan. 2024.